Spark Example on Quickstart VM - Using Real World Sample

This document is to run Spark on CDH.

Apache Spark has been receiving a lot of deserved attention lately. It is very understandable given the huge importance of distributed data processing for many companies and the pursuit for faster, cheaper and easier to use technologies aiming replace or complement the widely adopted Hadoop ecosystem and its MapReduce paradigm.

You need to some preparation to run this example.

1. VMware Download(http://www.vmware.com/kr) and Install
2. Quickstart tutorial Download(http://www.cloudera.com/content/cloudera/en/downloads.html)
3. Open the Quickstart tutorial file(/cloudera-quickstart-vm-5.3.0-0-vmware/cloudera-quickstart-vm-5.3.0-0-vmware.vmx) on VMware.
4. Download the dataset on Github.*

* Download dataset for running spark example.

Open the terminal on virtual machine of QuickStart tutorial. And input command.

```
[cloudera@quickstart ~] git clone https://github.com/SteveKim0513/spark_real_world_dataset.git
```

```
[cloudera@quickstart ~]$ git clone https://github.com/SteveKim0513/spark_real_world_dataset.git
Initialized empty Git repository in /home/cloudera/spark_real_world_dataset/.git/
remote: Counting objects: 16, done.
remote: Compressing objects: 100% (12/12), done.
remote: Total 16 (delta 2), reused 13 (delta 2)
Unpacking objects: 100% (16/16), done.
[cloudera@quickstart ~]$ ls
cloudera-manager  datasets  Documents  eclipse  Music    Public                 Templates  workspace
cm_api.sh         Desktop   Downloads  lib      Pictures spark_real_world_dataset Videos
[cloudera@quickstart ~]$
```

You can see the new folder named spark_real_world_dataset

```
[cloudera@quickstart ~] ls ./spark_real_world_dataset
```

```
[cloudera@quickstart ~]$ ls ./spark_real_world_dataset
life_expectation.csv  number_of_hospital.csv  property_tax.csv
local_tax.csv         park_area.csv           README.md
[cloudera@quickstart ~]$
```

These datasets are downloaded from “http://data.seoul.go.kr”.

< Scenario >

Life expectancy is the most commonly used measure to describe population health. Life expectancy measures how long, on average, a person is expected to live based on current age and sex-specific death rates. It is often expressed as the number of years of life a person born today is expected to live.

The starting point for calculating life expectancy is the age-specific death rates of the population members. If a large amount of data is available, the age-specific death rates can be simply taken as the mortality rates actually experienced at each age.

But life expectancy is not a simple. There are various factors to affect the result. For example, medical treatment level, dietary life, crime rate, etc. So I wonder that how well the value of life expectancy describes real world.

<Assumption>

Wealth, the condition of medical treatment, surrounding environment(e,g. green area ratio) may affect the result of life expectancy.

<How to measure>

Analyze between independent variables and the dependent variable using linear regression on Spark.

<Implementation>

Step 1. Upload the data to HDFS

```
//make the folder at hadoop file system
[cloudera@quickstart ~] sudo -u hdfs hadoop fs -mkdir /user/spark_example

//change the authority of this folder.
[cloudera@quickstart ~] sudo -u hdfs hadoop fs -chmod +rw /user/spark_example

//copy the local dataset to hadoop file system
[cloudera@quickstart ~] hadoop fs -copyFromLocal ./spark_real_world_dataset/ /spark_example

//check the files
[cloudera@quickstart ~] hadoop fs -ls /user/spark_example
```

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -mkdir /user/spark_example
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -chmod +rw /user/spark_example
[cloudera@quickstart ~]$ hadoop fs -copyFromLocal ./spark_real_world_dataset/*.* /user/spark_example
[cloudera@quickstart ~]$ hadoop fs -ls /user/spark_example
Found 6 items
-rw-r--r--   1 cloudera supergroup        477 2015-02-21 07:47 /user/spark_example/README.md
-rw-r--r--   1 cloudera supergroup       7549 2015-02-21 07:47 /user/spark_example/life_expectation.
csv
-rw-r--r--   1 cloudera supergroup      10114 2015-02-21 07:47 /user/spark_example/local_tax.csv
-rw-r--r--   1 cloudera supergroup       6627 2015-02-21 07:47 /user/spark_example/number_of_hospita
l.csv
-rw-r--r--   1 cloudera supergroup       7748 2015-02-21 07:47 /user/spark_example/park_area.csv
-rw-r--r--   1 cloudera supergroup       5752 2015-02-21 07:47 /user/spark_example/property_tax.csv
[cloudera@quickstart ~]$
```

Step 2. Refine the data using Hive

1) Open the Hue to run Hive on GUI environment

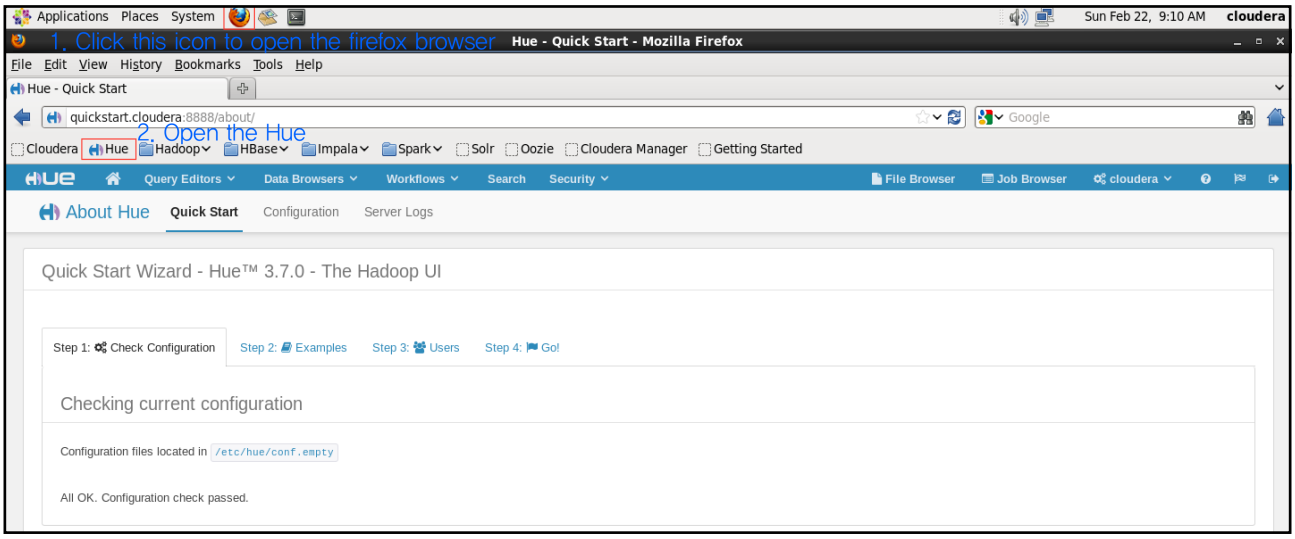

2) Click the meatball to create the table

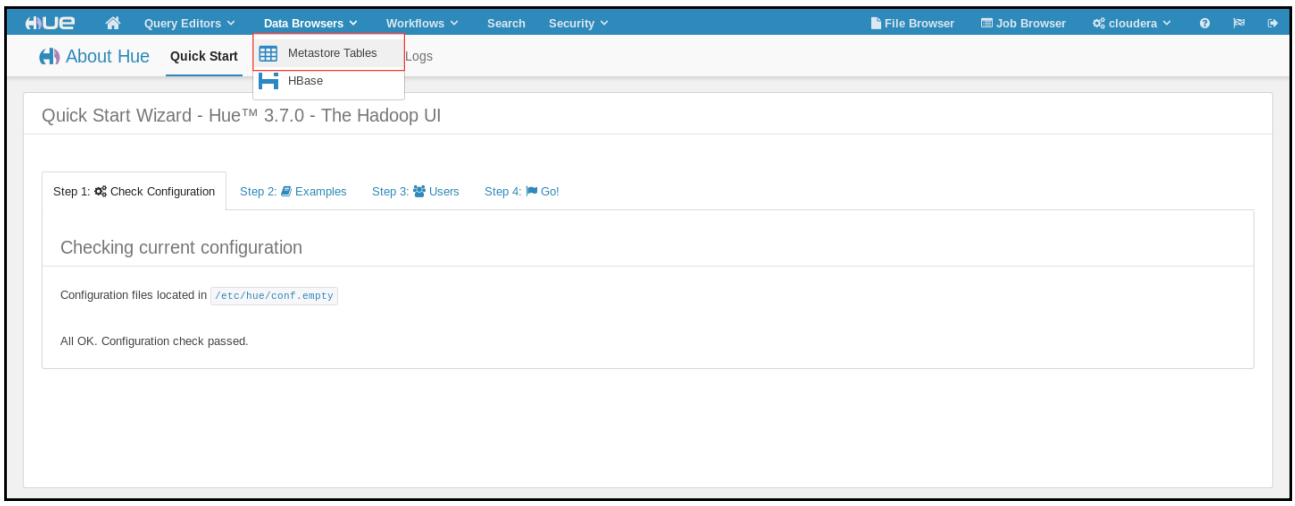

3) Create the table from a data file(life_expectation)

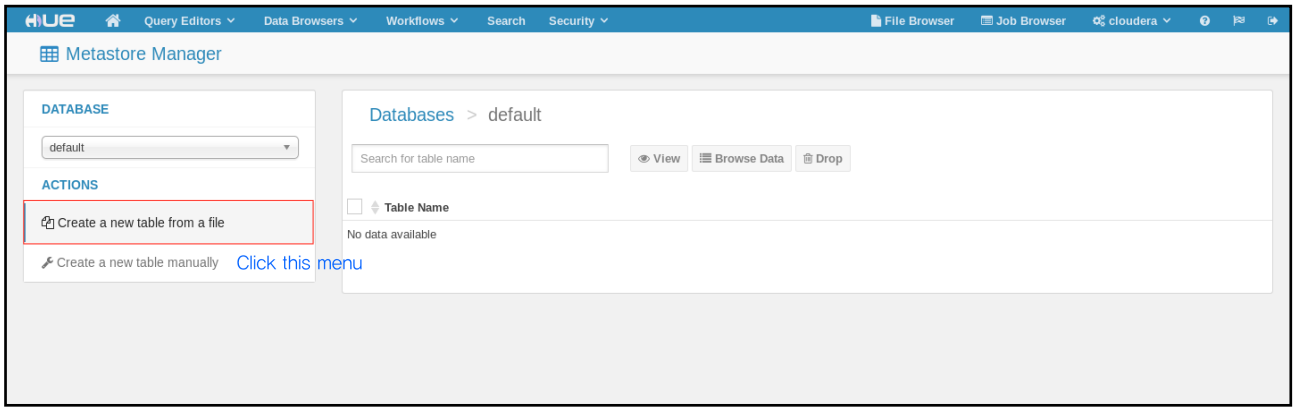

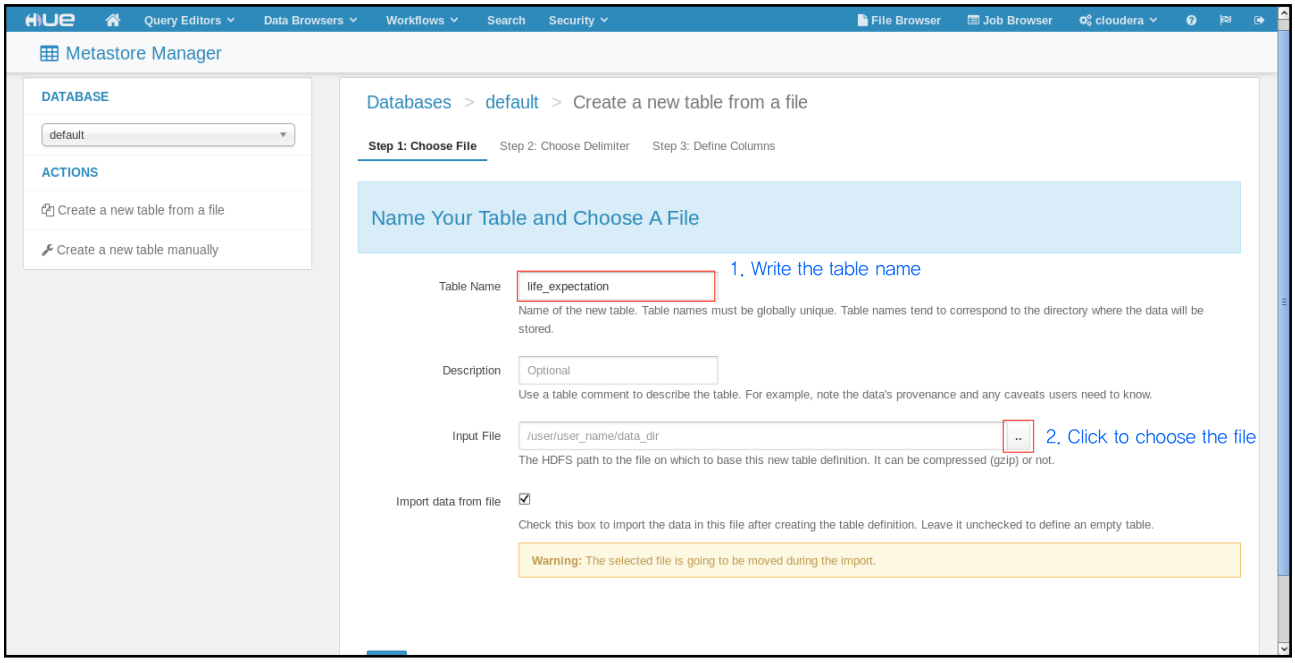

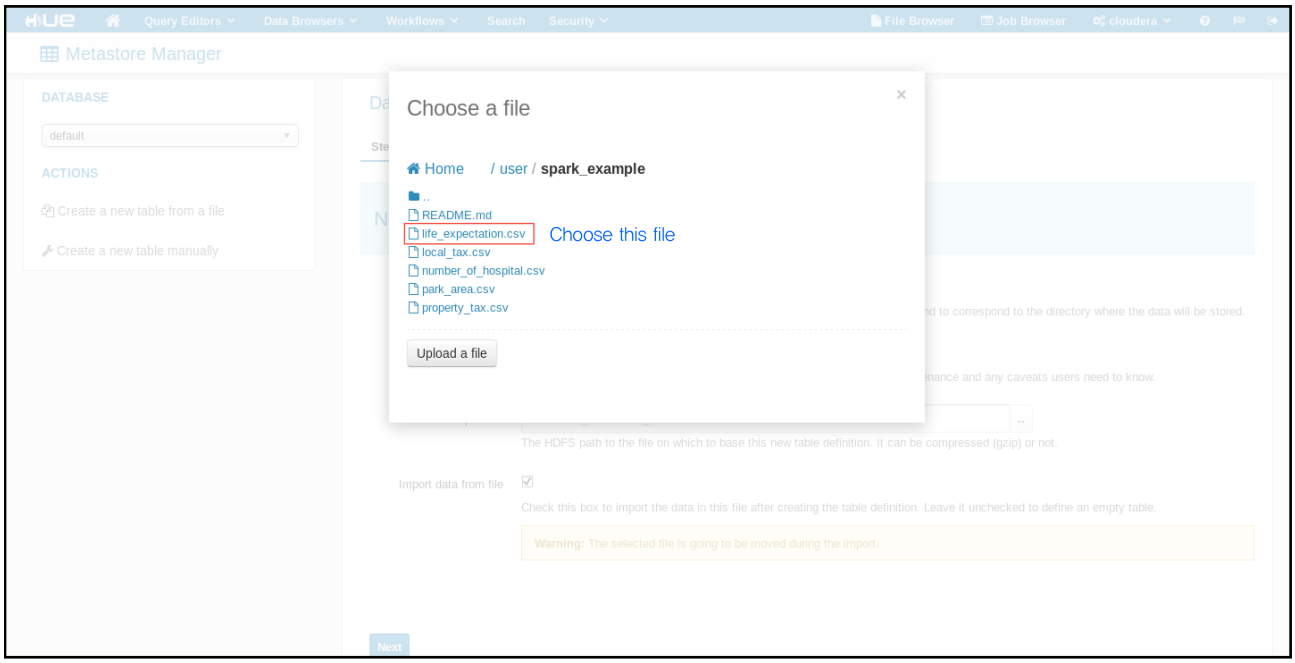
Metastore Manager
Choose a file
Home / user / spark_example
README.md
life_expectation.csv
Choose this file
local_tax.csv
number_of_hospital.csv
park_area.csv
property_tax.csv
Upload a file

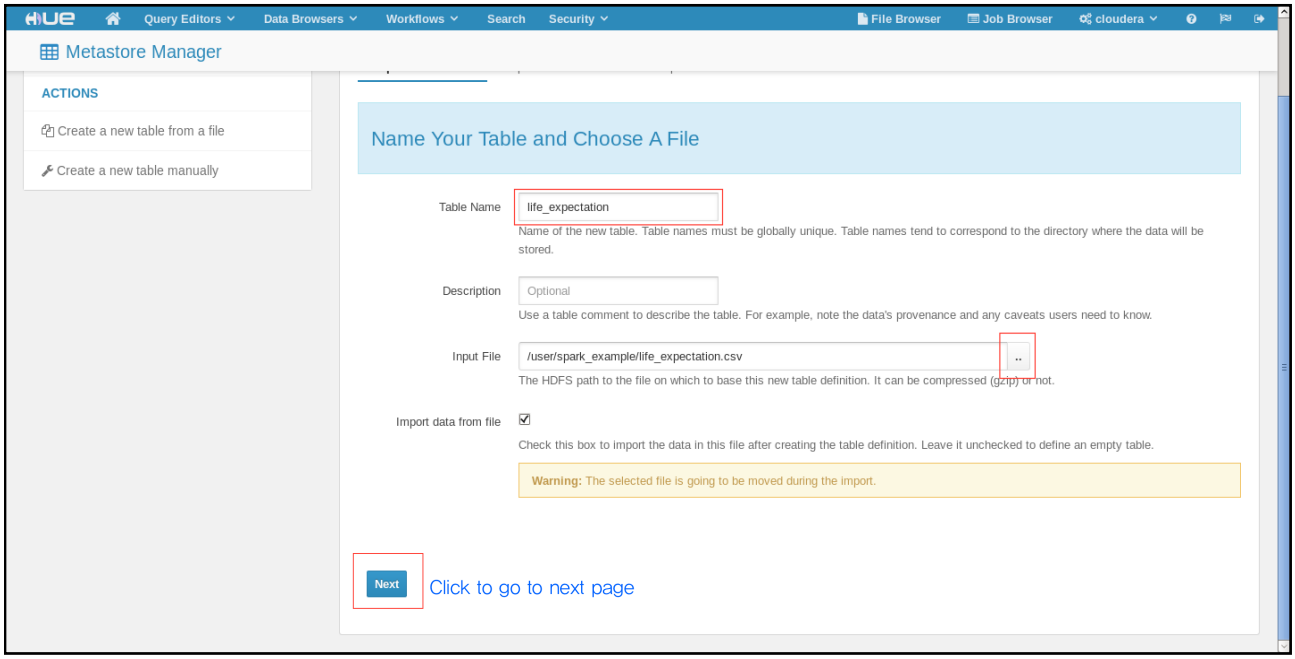
Metastore Manager
ACTIONS
Create a new table from a file
Create a new table manually
Name Your Table and Choose A File
Table Name
life_expectation
Name of the new table. Table names must be globally unique. Table names tend to correspond to the directory where the data will be stored.
Description
Optional
Use a table comment to describe the table. For example, note the data's provenance and any caveats users need to know.
Input File
/user/spark_example/life_expectation.csv
The HDFS path to the file on which to base this new table definition. It can be compressed (gzip) or not.
Import data from file
Check this box to import the data in this file after creating the table definition. Leave it unchecked to define an empty table.
Warning: The selected file is going to be moved during the import.
Next
Click to go to next page

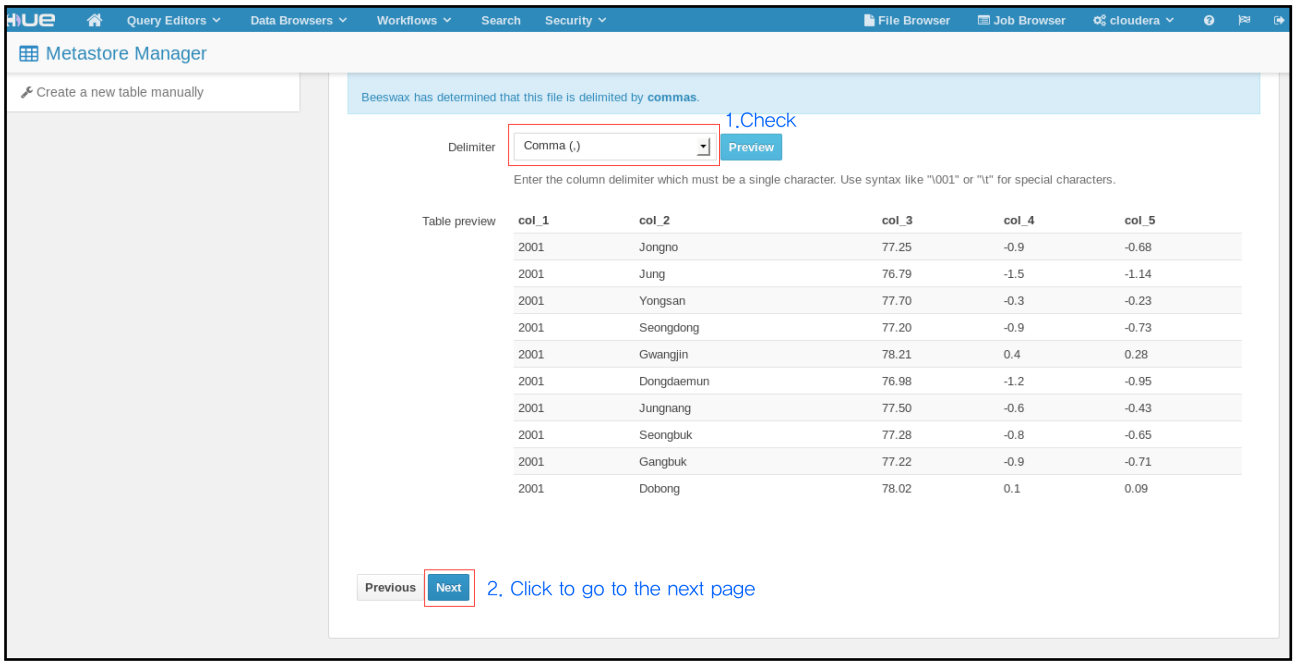
Query Editors
Data Browsers
Workflows
Search
Security
File Browser
Job Browser
cloudera
Metastore Manager
Create a new table manually
Beeswax has determined that this file is delimited by commas.
1.Check
Delimiter
Comma (,)
Preview
Enter the column delimiter which must be a single character. Use syntax like "\001" or "\t" for special characters.
Table preview
col_1 col_2 col_3 col_4 col_5
2001 Jongno 77.25 -0.9 -0.68
2001 Jung 76.79 -1.5 -1.14
2001 Yongsan 77.70 -0.3 -0.23
2001 Seongdong 77.20 -0.9 -0.73
2001 Gwangjin 78.21 0.4 0.28
2001 Dongdaemun 76.98 -1.2 -0.95
2001 Jungnang 77.50 -0.6 -0.43
2001 Seongbuk 77.28 -0.8 -0.65
2001 Gangbuk 77.22 -0.9 -0.71
2001 Dobong 78.02 0.1 0.09
Previous
Next
2. Click to go to the next page

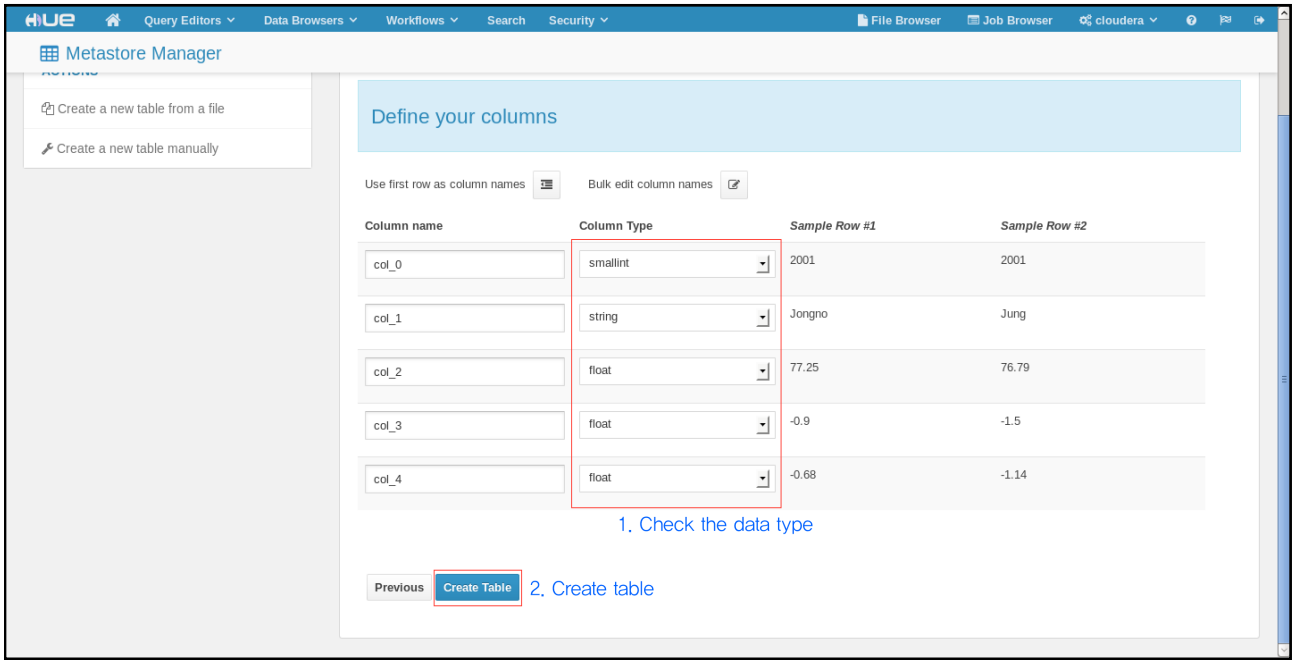
Query Editors
Data Browsers
Workflows
Search
Security
File Browser
Job Browser
cloudera
Metastore Manager
Create a new table from a file
Create a new table manually
Define your columns
Use first row as column names
Bulk edit column names
Column name
Column Type
Sample Row #1
Sample Row #2
col_0 smallint 2001 2001
col_1 string Jongno Jung
col_2 float 77.25 76.79
col_3 float -0.9 -1.5
col_4 float -0.68 -1.14
1. Check the data type
Previous
Create Table
2. Create table

4) Check the created table

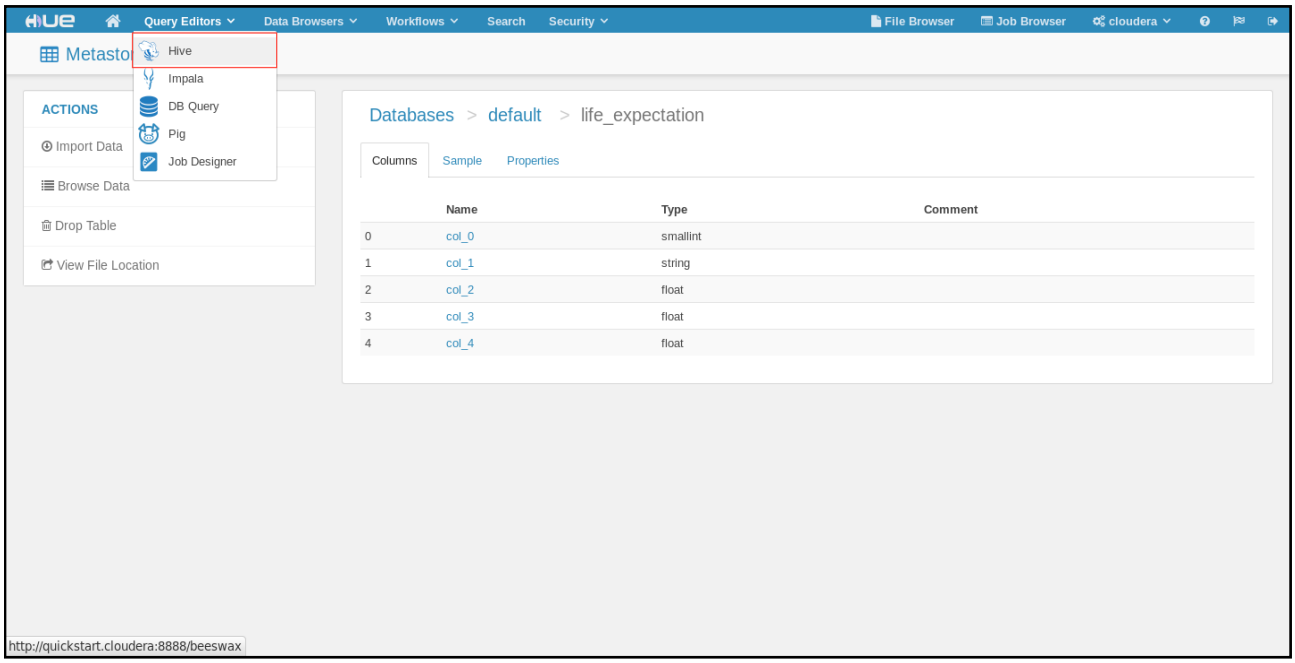

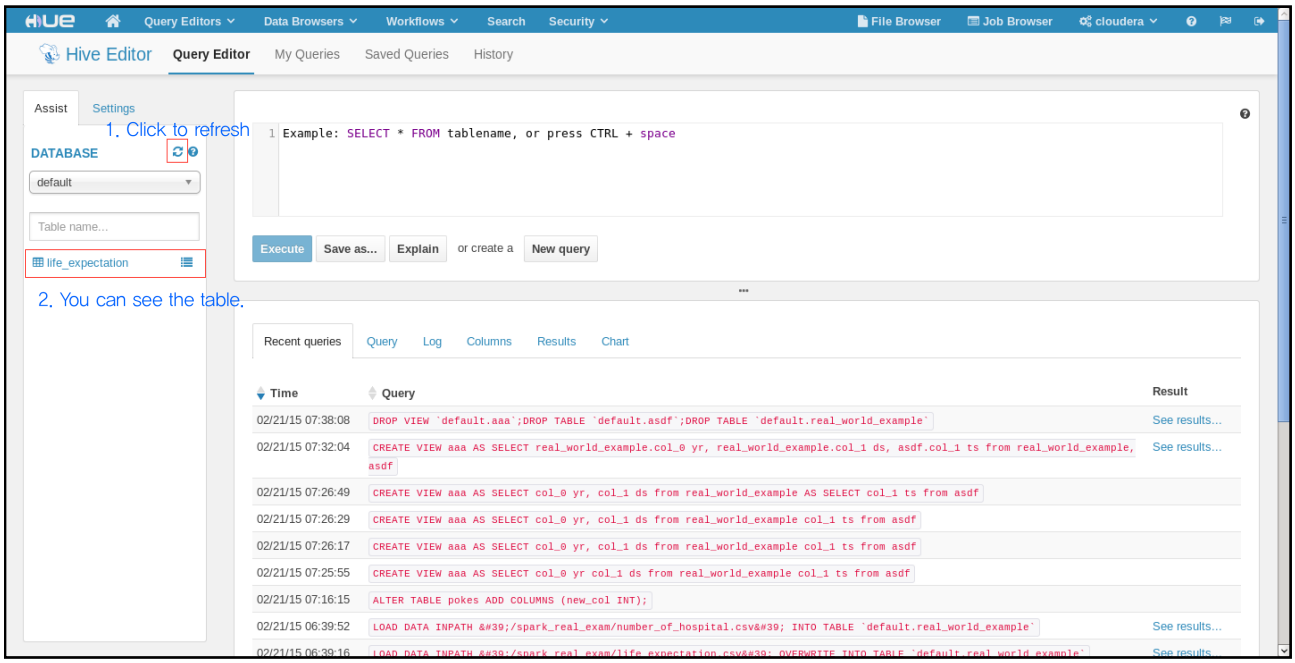

You have to create other tables using uploaded files. You can create the tables on the same way. And you must check the data type and change it if needed.

*cf. Change data type

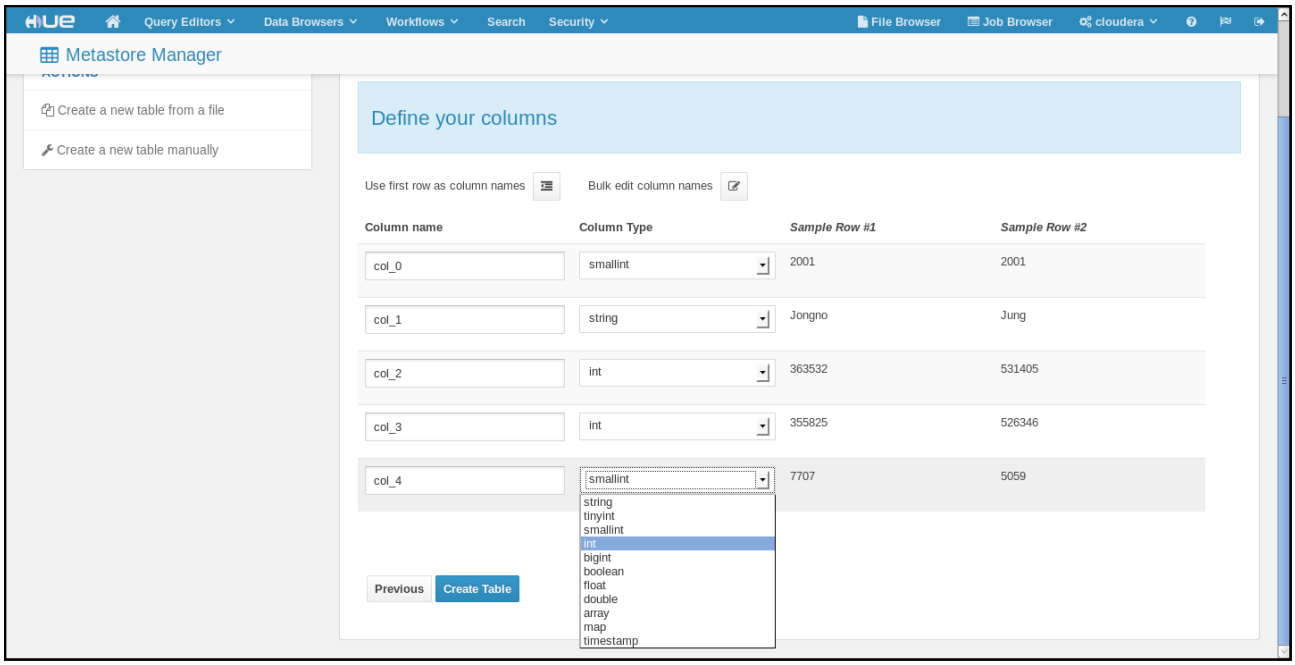

*cf. Final data type

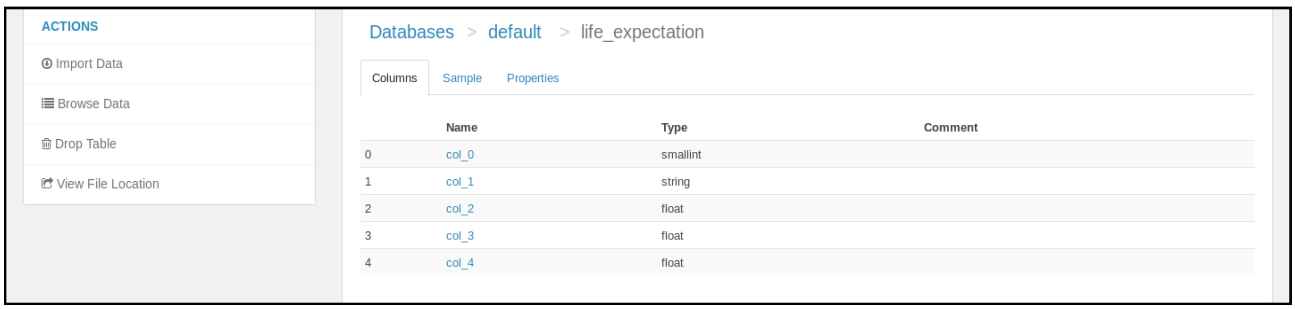

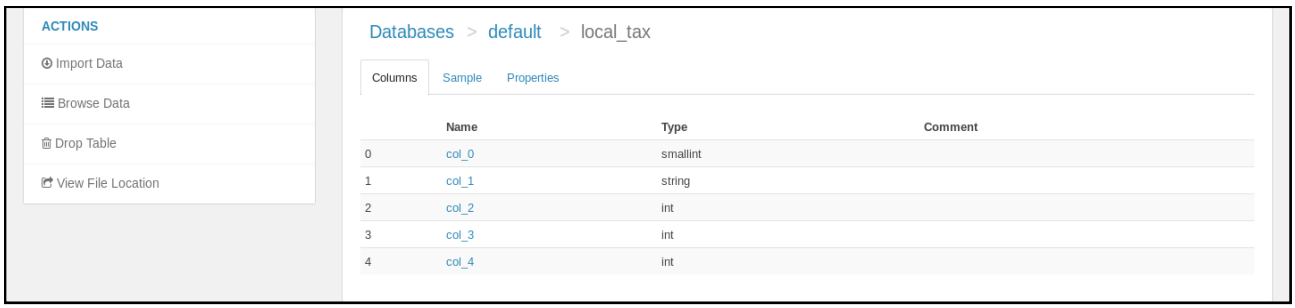

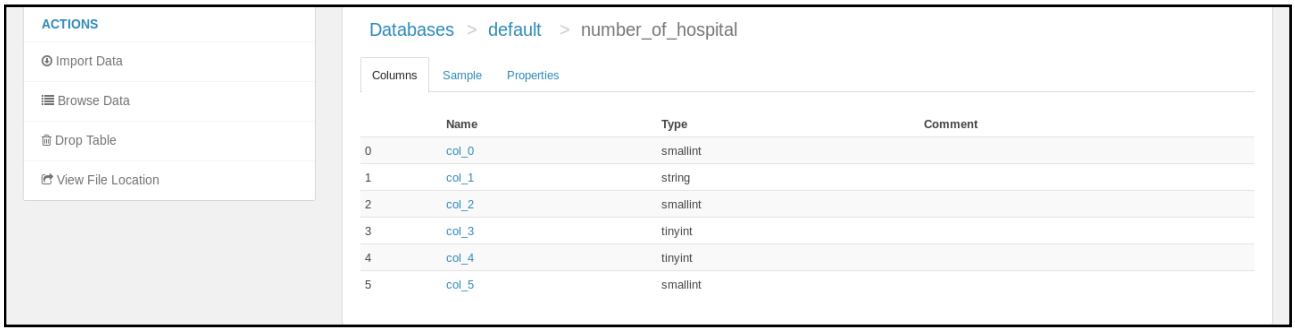

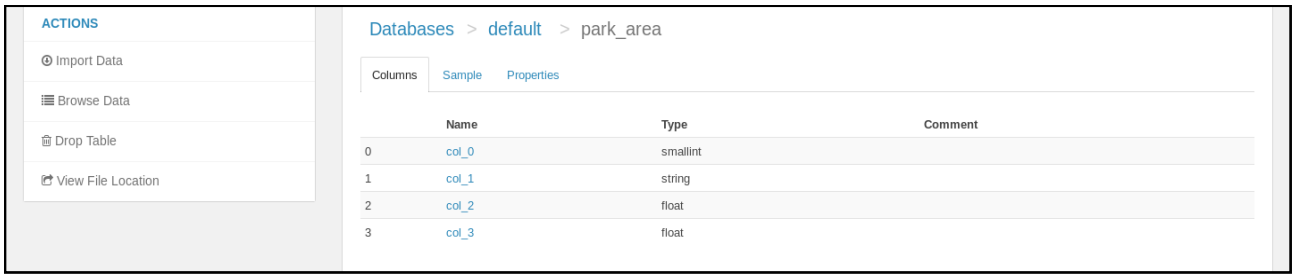

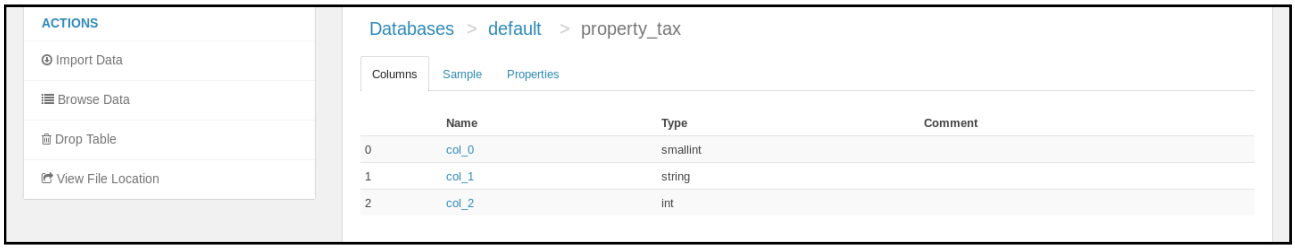

5) Result of creating tables

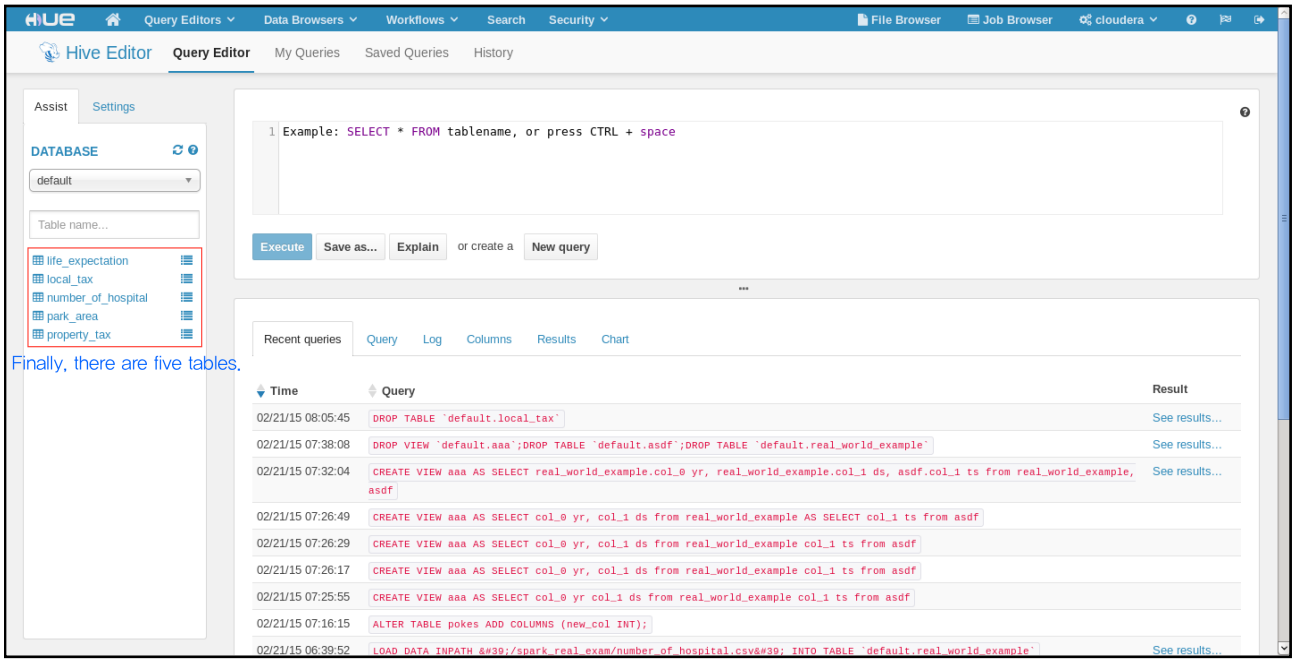

Finally, there are five tables.

6) Create view to refine data

Five data tables have a lot of data and some of them are not needed. So you have gather the needed data into one file.

```
CREATE VIEW linear_data AS SELECT
life_expectation.col_2 exp_yr,
local_tax.col_2 l_tax,
number_of_hospital.col_2 num_hospital,
park_area.col_3 per_area,
property_tax.col_2 p_tax
FROM life_expectation
JOIN local_tax
ON (life_expectation.col_0 = local_tax.col_0) AND (life_expectation.col_1 = local_tax.col_1)
JOIN number_of_hospital
ON (life_expectation.col_0 = number_of_hospital.col_0) AND (life_expectation.col_1 =
number_of_hospital.col_1)
JOIN park_area
ON (life_expectation.col_0 = park_area.col_0) AND (life_expectation.col_1 = park_area.col_1)
JOIN property_tax
ON (life_expectation.col_0 = property_tax.col_0) AND (life_expectation.col_1 = property_tax.col_1)
```

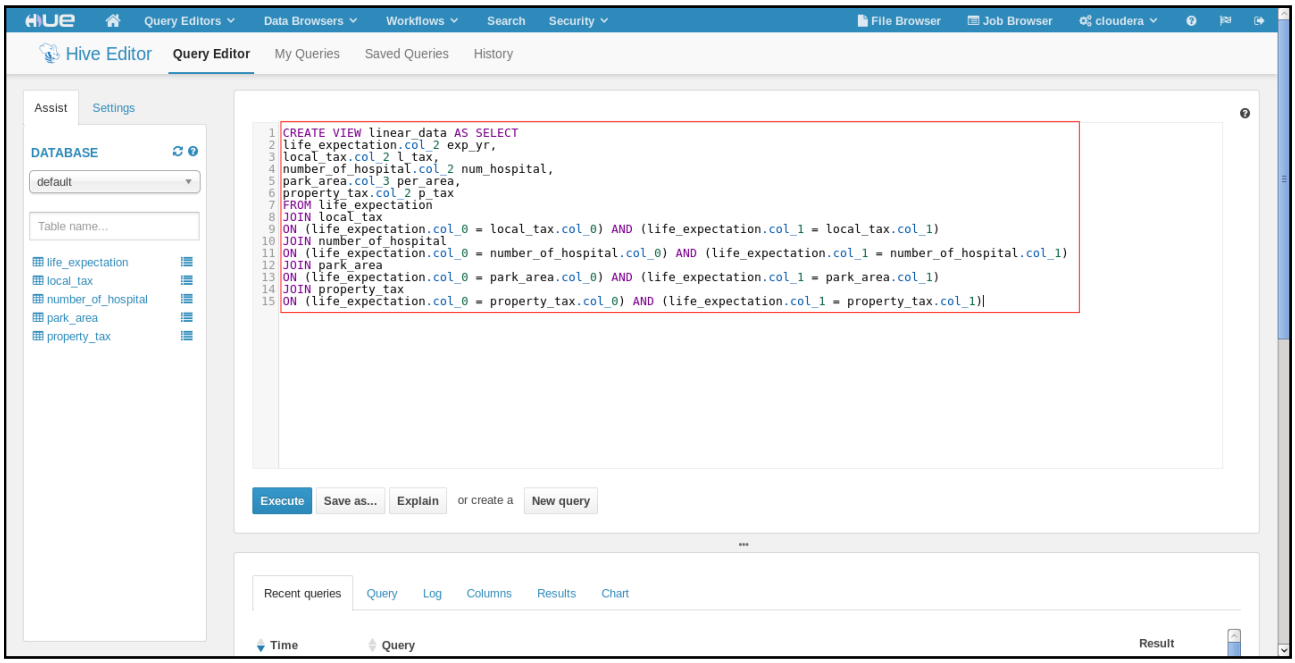

Click the “Excute" button.

7) Check the result of creating view

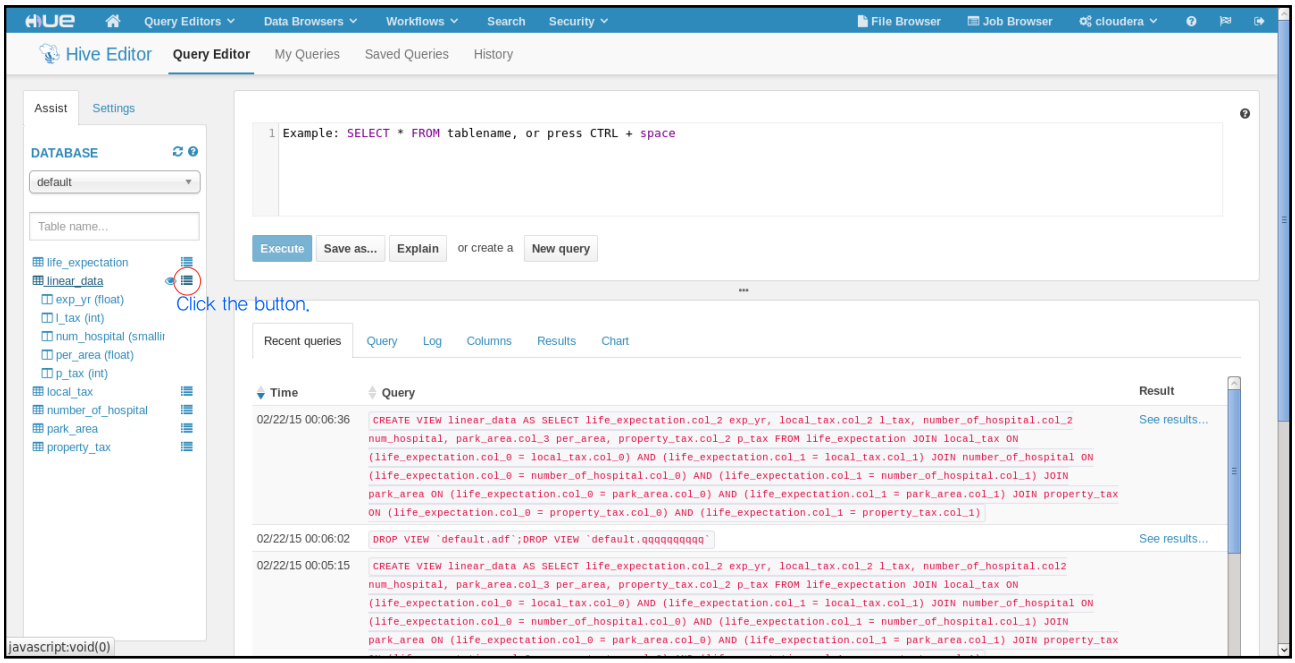

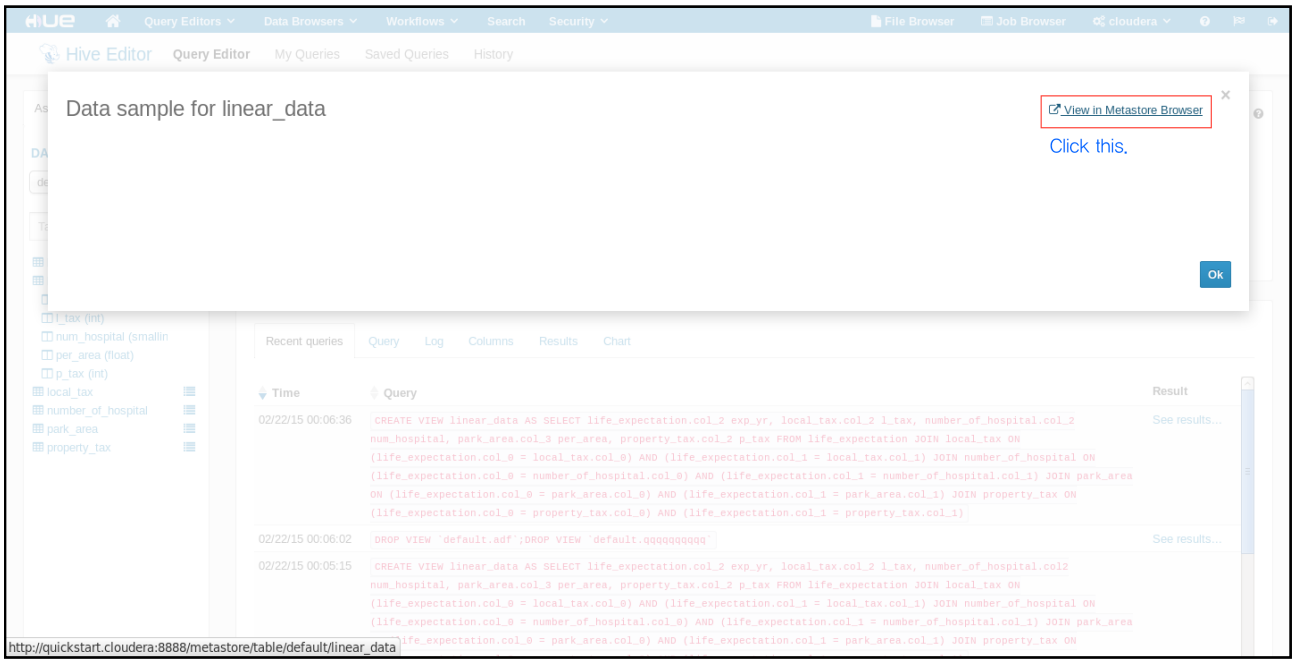

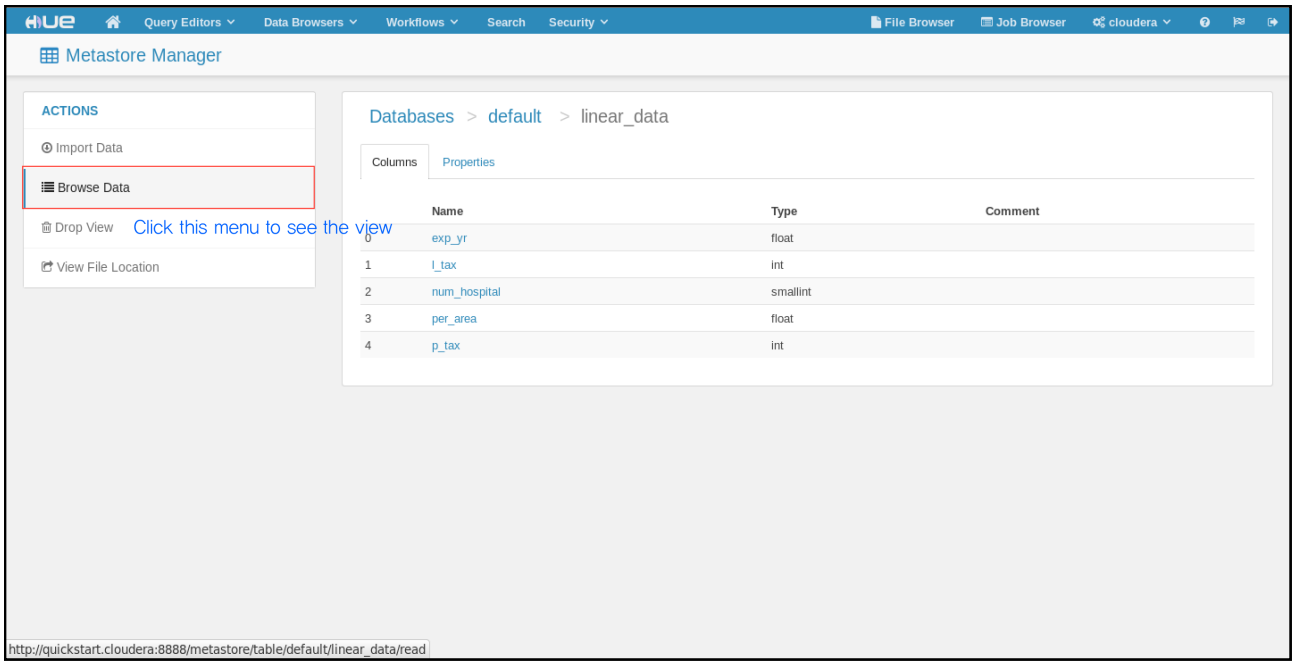

Now, you can see the refined dataset. There are five columns.
Column 1. Life expectation (y data)
Column 2. Local tax (x1 data)
Column 3. Number of hospital (x2 data)
Column4. Park area per a person (x3 data)
Column5. Property tax (x4 data)

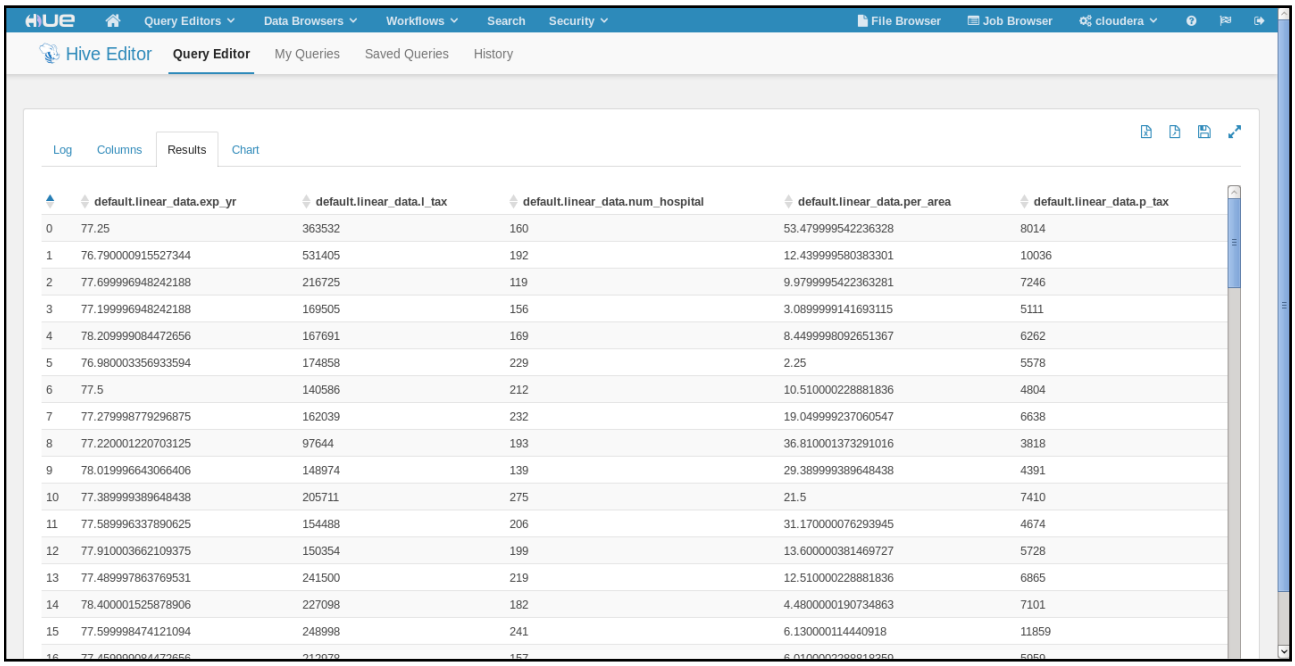

	default.linear_data.exp_yr	default.linear_data.l_tax	default.linear_data.num_hospital	default.linear_data.per_area	default.linear_data.p_tax
0	77.25	363532	160	53.479999542236328	8014
1	76.790000915527344	531405	192	12.439999580383301	10036
2	77.699996948242188	216725	119	9.9799995422363281	7246
3	77.199996948242188	169505	156	3.0899999141693115	5111
4	78.209999084472656	167691	169	8.4499998092651367	6262
5	76.980003356933594	174858	229	2.25	5578
6	77.5	140586	212	10.510000228881836	4804
7	77.279998779296875	162039	232	19.049999237060547	6638
8	77.220001220703125	97644	193	36.810001373291016	3818
9	78.019996643066406	148974	139	29.389999389648438	4391
10	77.389999389648438	205711	275	21.5	7410
11	77.589996337890625	154488	206	31.170000076293945	4674
12	77.910003662109375	150354	199	13.600000381469727	5728
13	77.489997863769531	241500	219	12.510000228881836	6865
14	78.400001525878906	227098	182	4.4800000190734863	7101
15	77.599998474121094	248998	241	6.1300001144409180	11859

8) Save the refined data to file on HDFS

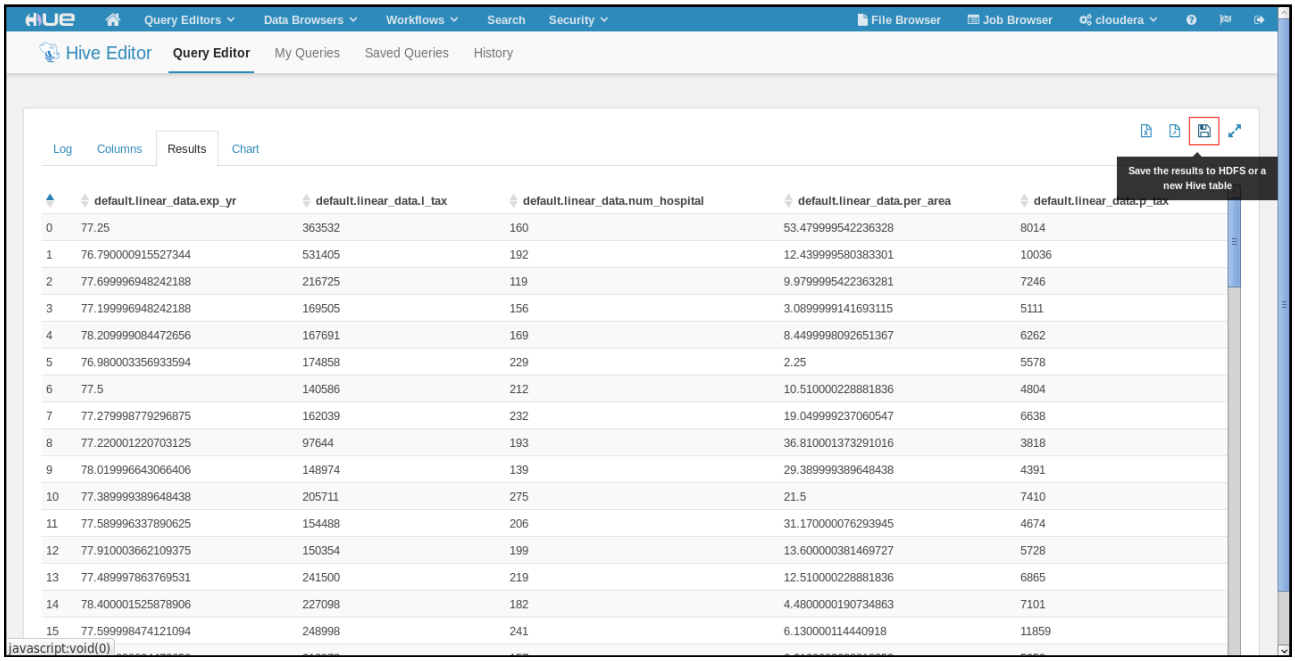

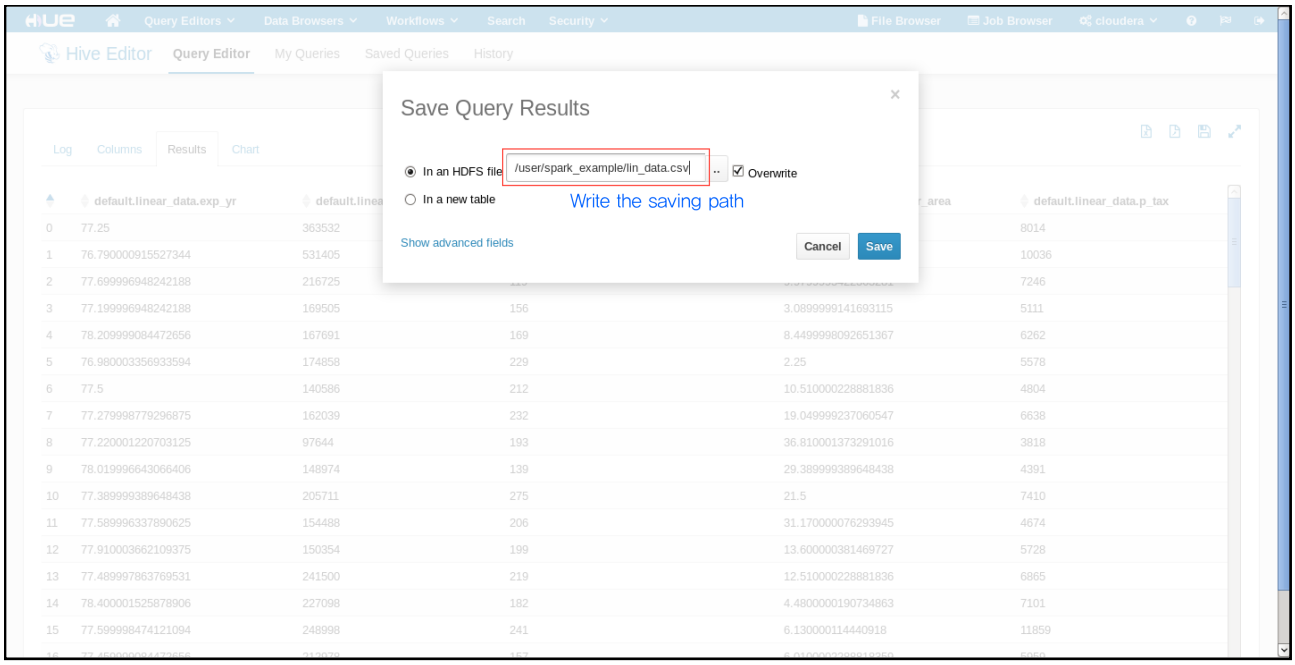

9) Check the refined file on HDFS

```
[cloudera@quickstart ~] hadoop fs -ls /user/spark_example
```

```
[cloudera@quickstart ~]$ hadoop fs -ls /user/spark_example
Found 2 items
-rw-r--r--   1 cloudera supergroup        477 2015-02-21 23:34 /user/spark_example/README.md
-rw-r--r--   1 cloudera supergroup      11283 2015-02-22 00:12 /user/spark_example/lin_data.csv
[cloudera@quickstart ~]$
```